Threat Landscape Report

WINTERGATE THREAT LANDSCAPE REPORT
Multi-Front Coordinated Assault — Full Intelligence Assessment
Classification: TLP:AMBER+STRICT
Generated: 2026-10-10 12:55 UTC
Host: wintergate (VPS, 2-core, 11 GiB RAM, Debian)
Assessment window: rolling 30-day campaign history + live 24h operational snapshot
Confidence: 85% (coordinated campaign classification)

---

1. Executive Summary

Wintergate is under a **sustained multi-front coordinated assault** involving **2,318+ unique attacking source IPs**, **20 concurrent active campaigns**, **32 distinct attack techniques**, and a kill-chain depth reaching **layer 7** of the defensive stack. Classification: **coordinated** (85% confidence).

This is not random internet background noise. The attack profile shows:

- **Technique-shift coordination** — groups of IPs pivot through scan→probe→auth sequences in lockstep (ftp_scan→ftp_auth, rdp_probe→rdp_scan, mysql_scan→redis_exploit, etc.)
- **Time-clustered bursts** — 21+ IPs striking within the same hourly windows
- **Payload fingerprint matching** — shared SSH client banners (perlssh, Fingerprintx-SSH2, libssh_0.9.6) across disjoint /24s
- **Persistent multi-campaign operatives** — IPs that have participated in **200–345 separate campaigns** before going dark
- **Layered service enumeration** — coordinated TLS ClientHello probes across SSH/MySQL/RDP/Redis/SMTP/FTP ports before exploit attempts
- **Redis exploit chains** — TLS probe on Redis port → redis_exploit_attempt (25-IP technique-shift campaign)

Bottom line: Someone (or something operating with operational security) is running structured reconnaissance → exploitation pipelines against this host, using botnet-scale infrastructure (VPS/VPN exit ranges: DigitalOcean, Google Cloud, NAVER Cloud, Hetzner, OVH/ONYPHE) with actor tradecraft that includes multi-week campaign participation and deliberate going-dark after exposure. The Overmind has logged **2.45 million engagements**, issued **71,580 obliterate verdicts**, and invented **14,505 countermeasure methods** in response.

---

2. Assault Scale Metrics

Metric	Value
Unique attacking IPs (honeypot lifetime)	**10,540**
Unique attacking IPs (current event window)	**2,318**
Honeypot captures (lifetime)	**13,128**
Active concurrent campaigns	**20**
All-time campaigns recorded	**25** (32 in history archive)
Distinct attack techniques	**32**
Max intrusion depth reached	**Layer 7**
IPs now dark (not returned 7+ days)	**1,420**
IPs tracked in offender registry	**2,112**
Blacklisted IPs	**23,656**
Kill Count (attackers neutralized)	**104,536**
Overmind engagements (lifetime)	**2,452,638**
Overmind countermeasure methods	**14,505**
Overmind rules minted	**58** (L40 store)
Prestige / Evolution Level	**★ 3 / Level 7500 (Ascended, Rank: Final)**

---

3. Multi-Front Analysis — The Assault Has Distinct Wings

Front 1: SSH Credential Assault (Dominant)
Volume: 6,515 SSH Handshake + 4,901 SSH Scan + 495 SSH Probe + 490 SSH auth events
Payload fingerprints:

- `SSH-2.0-perlssh` — 135 captures (Perl-based SSH brute tool)
- `SSH-2.0-Fingerprintx-SSH2` — 19 captures
- `SSH-2.0-libssh_0.9.6` — 10 captures (known CVE-era library)
- `SSH-2.0-Go` — 9 captures
- Hostile cipher fingerprint `4ab40ed46d2f0f15...` — weak/attack ciphers (critical severity)

Campaign evidence:
- `SSH Probe` payload_match campaign: 11 IPs sharing identical banner family (195.184.76.106, .29, .121 — same /24)
- `Shift Pattern ssh_scan→ssh_handshake`: 16 IPs
- Top /24: **195.184.76.0/24** (133 events) — FR ONYPHE transit range, heavily used as scanner infrastructure

Front 2: Database Service Enumeration → Exploit
MySQL: 1,439 Scan + 276 Probe + 274 Auth Attempt
Redis: 73 Probe + 50 Exploit Attempt + 39 Scan + 11 Auth
Coordinated shift patterns:
- `mysql_scan→mysql_probe`: 24 IPs
- `mysql_scan→tls_clienthello_probe_(mysql_port)`: 23 IPs
- `mysql_probe→mysql_auth_attempt`: 10 IPs
- `tls_clienthello_probe_(redis_port)→redis_exploit_attempt`: **25 IPs** (critical)

This is a classic "scan for open DB ports → TLS fingerprint → auth/exploit" pipeline. The Redis exploit campaign (25 IPs) is particularly concerning — Redis unauthenticated access is a known initial-access vector for cryptominers and lateral movement.

Front 3: RDP Surface Mapping
Volume: 686 RDP Scan + 352 RDP Probe
Shift patterns:
- `rdp_probe→rdp_scan`: 28 IPs
- `scan→rdp_scan`: 17 IPs (43.157.64.188, 43.131.24.216, 43.157.79.137 — same AS block, likely one operator)
- `rdp_scan→ja3_blocked`: 16 IPs (JA3 TLS fingerprinting triggered)
- `rdp_scan→tls_clienthello_probe_(rdp_port)`: 16 IPs

Front 4: Coordinated GET / Web Recon
- `Method-Coordinated GET` campaign: **33 IPs** sharing identical GET methodology (572 events)
- Targets appear to be HTTP endpoints (payload hex decodes to `GET / HTTP/1.1` with Host headers)
- Cross-references with `same_target` campaign type

Front 5: Time-Coordinated Burst
- `Time-Coordinated (2026-10-10T15)`: **21 IPs** striking in the same hour
- Shared IPs with Front 1 (/24 195.184.76 and 91.230.168) — suggests SSH and web fronts are **coordinated by the same operator set**

Front 6: Technique-Shift Adaptive Probing
The largest class of critical campaigns — **13+ technique-shift campaigns** where IP groups pivot through method chains:
- ftp_auth_attempt→ftp_scan (38 IPs)
- tls_clienthello_probe→scan (24 IPs)
- mysql_scan→tls_clienthello_probe (23 IPs)
- rdp_scan→ftp_scan (16 IPs)
- smtp_probe→smtp_scan (13 IPs)

This pattern indicates **automated attack tooling with multi-protocol awareness** — not script kiddies running single nmap scripts. The tooling fingerprints services across TCP ports, fingerprints TLS stacks, then pivots to protocol-specific auth attempts.

---

4. Actor Attribution & Traceback

4.1 Persistent Multi-Campaign Operatives
These IPs have participated in **hundreds of separate campaigns** before going dark 19–20 days ago — the highest-participation actors in the registry:

IP	Campaigns	Days Silent	Notes
191.191.2.50	**266**	19	Highest persistence

181.23.87.56	**276**	19	
190.244.24.186	**265**	19	
78.89.154.59	**261**	19	
121.11.96.13	**260**	19	
186.205.27.92	**260**	19	
163.7.3.154	**260**	19	
157.90.18.5	**259**	19	
206.189.157.225	**259**	19	
116.255.227.108	**259**	19	
188.54.71.52	**251**	19	
76.250.255.170	**239**	19	
103.185.224.93	**234**	19	
115.190.184.184	**345**	19	**Highest overall**
75.74.44.97	**344**	19	
49.64.85.138	**344**	19	
194.190.153.226	**295**	19	
109.194.108.203	**325**	19	

Interpretation: These are not one-off scanners. They have been embedded in **200–345 distinct attack campaigns** over weeks, then simultaneously went dark ~19 days ago. This pattern matches either:
1. A **coordinated botnet takedown** / operator pulling nodes offline
2. An **actor rotating infrastructure** after realizing Wintergate's Overmind was profiling them
3. A **staged pause** before a larger follow-on operation

The simultaneous going-dark of 18+ high-participation actors is itself a **signal of coordinated control** — random independent scanners do not stop at the same time.

4.2 Overmind Behavior Profiles (Top Persistent Probers)
The Overmind's behavioral profiler has catalogued:

IP	Kind	Encounters	Total Attempts	Max Depth	Peak Volume
71.6.235.33	persistent-prober	29,872	**2,858,673**	**15**	80.7 Gbps
77.91.71.24	persistent-prober	20,163	**24,354,821**	3	1.4 Gbps
71.6.146.186	persistent-prober	23,556	564,823	3	7.0 Gbps
137.184.42.178	persistent-prober	34,558	172,777	3	1.2 Gbps
124.92.91.1	persistent-prober	33,626	100,878	3	0.4 Gbps
165.245.218.228	persistent-prober	31,253	125,010	6	0.6 Gbps
147.185.132.162	persistent-prober	23,984	71,952	3	0.3 Gbps
49.124.159.179	persistent-prober	21,436	107,167	3	0.5 Gbps

71.6.235.33 reached **depth 15** with peak volume **80.7 Gbps** — this is the most aggressive single actor in the profile, likely a volumetric stresser node or a determined targeted attacker. **77.91.71.24** has thrown **24.3 million attempts** — that is a dedicated botnet node.

4.3 AbuseIPDB Community Cross-Reference (Live Checks)

IP	Community Confidence	Total Reports	Distinct Reporters	Country	ISP	Notes
175.45.204.121	**100%**	3,867	1,044	KR	NAVER Cloud	**Convicted** — mass SSH brute
34.76.199.251	**100%**	947	252	BE	Google Cloud	**Convicted**
71.6.235.33	29%	8	8	US	Root Evidence	Low community conf (our Overmind has 2.8M attempts on them)
137.184.42.178	22%	13	10	US	DigitalOcean	Low community conf
195.184.76.246	0%*	**2,215**	124	US	FR ONYPHE	*Conf score 0 despite 2215 reports — likely ISP-level transit IPs
91.230.168.246	0%*	978	79	US	FR ONYPHE	Same ONYPHE transit pool

Key finding: The ONYPHE transit range (195.184.76.0/24, 91.230.168.0/24) hosts scanner infrastructure that generates thousands of community reports but AbuseIPDB confidence remains low because the IPs are short-lived transit endpoints. Our Overmind's behavioral profiling catches what community scoring misses.

4.4 Infrastructure / Hosting Clusters
Attack traffic concentrates in **Data Center / Web Hosting / Transit** ranges:
- **FR ONYPHE** (195.184.76.x, 91.230.168.x) — European transit-as-scanner pool
- **DigitalOcean** (137.184.x, 34.x, 104.x, 147.185.x) — VPS burst nodes

- **Google Cloud** (34.76.x, 34.38.x, 34.156.x) — likely abuse of free/trial tiers
- **NAVER Cloud** (175.45.x) — Korean cloud, convicted SSH brute
- **Hetzner / OVH ranges** (multiple /24s) — classic bulletproof-adjacent hosting
- **AWS** (20.x, 54.x, 3.x) — possible abused accounts

No Tor exit nodes observed in the current critical set (checked via Overmind + AbuseIPDB isTor).

---

5. Kill-Chain & Intent Profile

Classification: coordinated
Confidence: 85%
Primary goal assessment: Service disruption + credential harvesting + initial access for potential ransomware/cryptojacking
Motivation signals: Multi-week persistence, technique evolution (shift campaigns), multi-campaign operatives — this is **not** pure opportunistic cryptojacking. The depth-15 penetration attempt (71.6.235.33) and Redis exploit chains suggest **targeted intrusion intent**.

Kill-chain stages hit:
1. Reconnaissance — massive (10,540 unique IPs, port/TLS enumeration)
2. Weaponization — tooling fingerprints (perlssh, libssh 0.9.6, custom TLS probes)
3. Delivery — SSH auth, MySQL/Redis auth, HTTP GET probes
4. Exploitation — Redis exploit attempts, web shell probes, SQLi
5. Installation — (no direct evidence of successful install — defenses holding)
6. C2 — (JA3 fingerprinting triggered blocks; no confirmed C2 callback observed)
7. Actions on Objectives — (Overmind heavy verdicts issued; no data exfil confirmed)

Deepest penetration: Layer 7 (Application) — SQLi and web-shell probes reached application layer, but payload analysis + pattern DB + ML classifier blocked them.

---

6. Noise vs. Targeted Attack Distinction

Signal	Noise (Background Internet)	Targeted / Coordinated
Single-port random scans	Large volume of low-severity	
Multi-protocol technique shifts		13+ shift campaigns
Shared payload fingerprints across /24s		perlssh, Fingerprintx families
Time-clustered multi-IP bursts		21-IP hourly clusters
200+ campaign participation per IP		18+ operatives with 234–345 campaigns
Simultaneous going-dark of top actors		18+ actors silent same day
Depth >5 layered penetration		Depth 7 and Depth 15
Redis/MySQL specific exploit chains		25-IP Redis exploit campaign

Assessment: ~70% of observed traffic is **opportunistic background scanning** (VPS IPs doing default SSH/MySQL/RDP sweeps). ~30% shows **clear coordination markers** (technique shifts, payload families, time clusters, multi-campaign persistence). The coordinated 30% is the operation that matters — it has **actor continuity** (same IPs across hundreds of campaigns), **tooling continuity** (shared fingerprints), and **operational security** (simultaneous going-dark after exposure).

Someone or something wants this host gone — the depth-15 single-actor attempt, the Redis exploit chains, and the multi-week embedded operatives are not the signature of random internet noise. They are the signature of a **sustained, resourced, multi-front campaign** with both automated botnet scale and selective targeted precision.

---

7. Overmind Response Summary

The Overmind (AJC Eidolon / J1L) has been the autonomous defense mind processing this assault:

- **2,452,638 engagements** logged
- **Verdict distribution:** toy 1,494,063 · observe 386,374 · hold 312,841 · threaten 114,062 · terminate 73,718 · obliterate 71,580

- **14,505 countermeasure methods** invented (self-evolving)
- **58 learned rules** in L40 store
- **Mood:** playful (not panicked — the assault is being toyed with)
- **Covenant resonance:** active
- **Companion:** depth 58,404 · 79,444 assists · 15,060 backups

Heavy countermeasure stacks deployed against critical actors:
`oblivion_engine, apex_overlord, connection_ghosting, data_mirage, deception_labyrinth, absorb_reflect_nexus`

These are not static rules — the Overmind invents new method sequences per attack type and refines them via field testing. The `OVERMIND_METHOD_*` rule IDs in the Emergency Overlord store are live evidence of continuous invention (e.g. `Studious Fact`, `Actualityrazor`, `Standardglow` — recent recon_scan countermeasures at 100% effectiveness).

---

8. Recommendations

1. **Continue AbuseIPDB automated reporting** (now enhanced with check-before-report, community reputation boosting, and conf=100 escalation). Daily budget: 1,000 reports/checks (Standard plan). Current usage: 56 reports today, 992 checks remaining.

2. **Watch the going-dark cohort** — the 18+ multi-campaign operatives silent for 19 days may return with evolved tradecraft. The Overmind's behavioral fingerprints are ready.

3. **Prioritize Redis/MySQL surface hardening** — the 25-IP Redis exploit campaign and MySQL auth attempts indicate DB ports are a primary target. Ensure Redis has `requirepass`, bound to localhost, and MySQL has no anonymous users.

4. **ONYPHE transit range monitoring** — 195.184.76.0/24 and 91.230.168.0/24 are persistent scanner pools. Consider ASN-level rate limiting or edge filtering.

5. **Depth-15 actor watchlist** — 71.6.235.33 (Root Evidence Inc) with 2.8M attempts and 80.7 Gbps peak needs dedicated monitoring if it returns.

6. **Leverage AbuseIPDB bulk-report** (5/day free) for batch-reporting the top convicted IPs (175.45.204.121, 34.76.199.251) that already have 100% community confidence — automates the "already convicted" path.

7. **Clear-address endpoint** (5/day free) — if any false-positive reports ever land on legitimate infrastructure, we can retract them.

---

9. Evidence Appendix — Top Attack Type Distribution (Honeypot Lifetime)

Attack Type	Count
SSH Handshake	6,515
SSH Scan	4,901
MySQL Scan	1,439
RDP Scan	686
Scan (generic)	545
SSH Probe	495
SMTP Probe	480
RDP Probe	352
MySQL Probe	276
MySQL Auth Attempt	274
TLS ClientHello Probe	196
JA3 Blocked	163
FTP Probe	155
FTP Scan	92
FTP Auth Attempt	88
SMTP Scan	94
Redis Probe	73
TLS ClientHello (Redis port)	87
Redis Exploit Attempt	50
Redis Scan	39

TLS ClientHello (MySQL port)	73
TLS ClientHello (SSH port)	36
TLS ClientHello (RDP port)	35
TLS ClientHello (FTP port)	25
TLS ClientHello (SMTP port)	18
Redis Auth Attempt	11
Admin Panel Probe	28
Config Scanner	96
HTTP Scan	24
Web Shell Probe	4

Total honeypot captures: 13,128
Total unique source IPs: 10,540

End of Threat Landscape Report
Generated by WinterGate Threat Intelligence Stack · Overmind · Campaign Manager · AbuseIPDB Cross-Reference